Phase 1 Corpus Statistics Report

Information Retrieval Project · arXiv Hybrid Search · Category: cs.CL (Computation and Language)

400	4 × 100	2023–2025	2 – 144
papers	sub-topics, balanced	0 papers outside the range	pages (median 15)

1. Category and sub-topics

All papers come from arXiv category **cs.CL** and have cs.CL as their *primary* category. The corpus covers four distinct NLP tasks: **Machine Translation**, **Question Answering**, **Summarization** and **Dialogue Systems**. Each paper's sub-topic label is the query it was retrieved by, and no paper appears under two sub-topics.

2. Collection method and reproducibility

- **Source:** arXiv API (Python `arxiv` package), PDFs downloaded from arXiv. Text and page counts were extracted with `pypdf`.
- **Date range:** 2023-01-01 to 2025-12-31, split into 12 quarters. Each sub-topic gets 9 papers per quarter in 2023 and 8 per quarter in 2024–2025 (4×9 + 8×8 = 100).
- **Query pattern** (48 queries = 4 sub-topics × 12 quarters, all stored in `collection_info.json`): `cat:cs.CL AND <phrase> AND submittedDate:[YYYYMM010000 TO YYYYMMDD2359]` Phrases: `"machine translation"`, `"question answering"`, `summarization`, `"dialogue systems"`. Results are sorted by submission date, newest first.
- **Filters:** primary category is cs.CL; the arXiv ID's year-month falls inside the queried quarter; no duplicates across sub-topics; papers of exactly 6 pages excluded (6 fits none of the defined tiers); papers with under 1,000 characters of extractable text excluded. Papers whose PDF returned HTTP 404 were skipped and replaced by the next result.
- **Stored fields per paper:** title, full text, arXiv ID, arXiv link, sub-topic, page count, length tier, submission date, quarter. Everything is stored locally in `corpus.csv`.

3. Length tiers

Short = fewer than 6 pages · Medium = 7–9 pages · Long = more than 9 pages (page count of the full PDF, including references and appendices).

Sub-topic	Short	Medium	Long	Total
Dialogue Systems	8	21	71	100
Machine Translation	9	21	70	100
Question Answering	5	11	84	100
Summarization	4	13	83	100
All	**26 (6.5%)**	**66 (16.5%)**	**308 (77.0%)**	**400**

All three tiers are present in every sub-topic. Long papers dominate because most cs.CL arXiv papers are conference or journal versions that include references and appendices, which push them past 9

pages. Page counts range from 2 to 144, with a median of 15.

4. Distribution over time

Year	Papers	Per quarter (all sub-topics)
2023	144	36
2024	128	32
2025	128	32
Total	**400**	

A check on arXiv ID prefixes finds 0 papers outside 2301–2512. Because results are sorted newest first within each quarter, every quarter's papers come from that quarter's last month (IDs 2303, 2306, 2309, 2312, 2403, …, 2512).

5. Development note

A first version of the corpus was collected newest-first over the whole 2023–2025 window and capped at 100 papers per sub-topic. It covered only April 2025 to February 2026, and 17 papers had 2026 IDs despite the date filter. The corpus was therefore re-collected with the quarterly design above, which spreads the papers over the full range and leaves 0 papers outside it.